

Enterprise User Perceptions of Agent Memory

UX Research: Sarah Hanses

April/May 2026

AGENDA

- 1 OVERVIEW
- 2 EXECUTIVE SUMMARY
- 3 DETAILED FINDINGS
- 4 DESIGN OPPORTUNITIES

1

OVERVIEW

OVERVIEW

What is memory in AI systems?



Memory is the **ability of an AI system to retain and recall information** about users - including facts, preferences, conversation history, past actions, and data from connected tools - across interactions, sessions, or tasks, and to use that information to shape future behavior.

OVERVIEW

Why this matters



Traditional LLMs operate in a stateless paradigm. Each interaction exists in isolation, with no knowledge carried forward from previous conversations. While this approach works basic tasks and short-lived agents, it fundamentally limits what AI systems can achieve.

The shift from stateless LLMs to [stateful agents](#) represents an evolution towards systems that can actually learn and adapt over time.

OVERVIEW

Why this matters



Memory **amplifies both the benefits and the risks of AI systems.**

When implemented responsibly, it can build trust and enable continuity, personalization, and more capable assistance and agentic workflows.

When implemented poorly, it will create risk including privacy violations, sensitive-data misuse, context collapse, data exfiltration, and the entrenchment of inaccurate or outdated information.

OVERVIEW

Top 8 Emerging Insights on AI memory

AI Memory User Experience in Enterprise Settings: Top 8 Emerging Research Insights

User Research: Sarah Hanses, AIFI (AI Futures & Insights) Team

AIFI Team Research: aka.ms/AIFutures

Special Thanks: To all who have conducted research on this topic (resource list [here](#)) and to all who provided feedback on this presentation.

Last updated: March 26, 2026

A recent synthesis on what we already know from existing research on AI memory surfaced 8 insights that were observed repeatedly across from 24 internal research studies.

This synthesis also surfaced open, unanswered questions on this topic, which led to this research.



OVERVIEW

Research Objectives



- 1 Learning Models: Explicit vs. Implicit**
Learn under what conditions an agent should learn implicitly from repeated user behavior vs. require explicit instruction.
- 2 Memory Architecture & Lifecycle**
Explore how users expect memory lifecycle and organization of memory in management settings to work.
- 3 Collaboration & Shared Context**
Uncover user expectations for how memory should work in collaborative contexts.
- 4 Governance, Privacy & Risk**
Understand important needs for governance, privacy and control of agent memory in enterprise settings.
- 5 Memory Models: Human vs. Agent**
Understand how users expect AI memory to mimic or differ from human memory.

OVERVIEW

Methodology

**All Personas**

Session Type	Moderated 1:1 Interviews
Session Length	45-60 minutes
Participants	15 total, including: 4 IT Admins, 3 AI Manager, 4 Citizen Makers, 4 Information Workers
Location	Remote via Teams or UserTesting
Discussion Topics	1. Warm-up 2. Learning & Adaptation (Implicit vs. Explicit learning) 3. Memory in Collaborative Contexts 4. Memory Lifecycle and Organization 5. Governance, Privacy & Risk 6. Mental Models: Human vs. Agent Memory 7. Wrap-up

[Link to discussion guide](#)

OVERVIEW

Personas



IT ADMINS

IT Admins are technical professionals who **plan, manage, and secure the systems** their organizations rely on. They administer Microsoft M365, and often platforms like Azure/AAD, SharePoint, or Power Platform at the tenant or environment level.

They keep data secure, monitor performance, troubleshoot issues, and manage licensing and access.

[\(screening criteria\)](#)



AI MANAGERS

AI Managers are leaders who oversee **how AI is planned, deployed, and scaled** across their organization, without necessarily building the solutions themselves.

They identify opportunities to improve processes or performance, align stakeholders, and ensure the right systems, data, and infrastructure are in place so AI initiatives can succeed and grow.

[\(screening criteria\)](#)



CITIZEN MAKERS

Citizen Makers are knowledge workers **who build solutions**, whether or not it's their core job, to improve how work gets done. They spot inefficiencies in processes and create new, often software-driven ways to reduce them.

They are non-technical innovators embedded within business functions who are transforming how work gets done by integrating AI, agentic tools, and automation directly into their day-to-day workflows.

[\(screening criteria\)](#)



INFORMATION WORKERS

Information workers in large enterprises who rely heavily on computers to perform their day-to-day tasks.

They work full-time across a wide range of professional roles, from operations and finance to marketing, HR, legal, healthcare, and technology. They have an **enterprise Microsoft 365 license** and **M365 Copilot app installed**.

[\(screening criteria\)](#)

2

EXECUTIVE SUMMARY

Executive Summary

01

Learning Models: Explicit vs. Implicit

Users want a hybrid model, implicit learning for routine, low-risk work and explicit teaching where rules are stable or stakes are high.

02

Memory Architecture & Lifecycle

Organize memory management around work context with flexible views. Treat lifecycle as context-dependent and.

03

Memory in Shared Contexts

Users see clear value in shared memory for handoffs, faster onboarding, but only with clear boundaries for individual vs. shared memories.

04

Governance, Privacy & Risk

Privacy and oversight are non-negotiable; users expect a shared, role-based governance model with strict guardrails on sensitive data.

05

Memory Models: Human vs. Agent


Agent memory should outperform humans on factual retention rigor, intentional forgetting, and auditability. Human memory should be used for emotional and relational discretion.

3

Detailed Findings


Learning Models: Explicit vs. Implicit

What is explicit vs. implicit learning?




Copilot – memory updated notification

Explicit prompt, with system inferring importance



Claude.md memory file example

Explicit instructions, provided by user



Claude – auto memory

Implicit learning

Note: These examples represent a small set of ways memory has shown up



Users prefer a hybrid model: AI should learn implicitly for routine work, but explicitly when rules are stable or risk is high

Should agents learn workflows implicitly from repeated patterns?


Are there scenarios in which users prefer implicit learning by the AI versus wanting to explicitly instruct it?

The **preferred model is a blend of implicit learning and explicit teaching.**

Users want AI to reduce repetitive prompting while keeping them in charge.

- *Automatic learning is best for low-stakes, individual preferences, and repetitive work like routine tasks, formatting and tone.*
- *Explicit teaching is preferred for fixed information or high-risk areas such as compliance, policy, and stable rules where mistakes are costly.*

A strong middle ground is suggestion-based learning: users are comfortable when the system proposes what to remember, but they choose whether to save it and where it applies.



*We should teach when the mistakes are costly or the rules are stable.
We let the agents learn when the patterns are evolving are ambiguous. (P2)*

This balance is so hard to get down. There are so many times I talk to it and I think you should have known, I do it this way every time. But there are times when it assumes too much. (P8)

Preferred learning models by task and relative risk – in draft

Implicit learning preferred


Routine + low-risk work

- Repetitive operational tasks
- Formatting, tone, and recurring workflows
- Patterns the agent can pick up to save time

Explicit teaching preferred

Stable rules and deterministic mappings

- Compliance and policy requirements
- Brand standards
- Taxonomies, fixed labels, source-of-truth instructions



AI should remember based both on user requests and evidence that the memory is useful

How should an agent learn from experience?

What signals or triggers (e.g., repeated behaviors, user cues) should cause an agent to form a new memory?


Direct signals should always count most: commands like *always remember*, *save*, *moving forward*, with stated preferences and instructions.

Repeated behavior should trigger memory suggestions: recurring or routine requests, or repeated formatting/tone preferences, are strong cues that something may be worth storing.

Usefulness matters more than frequency alone: users want AI to remember what is consistently accepted, rarely corrected, and linked to good downstream outcomes.

Human edits, overrides, and QA review are key learning signals, especially for subjective, high-stakes, or regulated work where AI should not self-validate.

Memory should be scoped and managed over time: store what is stable, avoid overgeneralizing context-specific information, require stronger permission for sensitive data, and treat memories as temporary until validated.



User requests

Usefulness evidence

I think if you tell it "always use this tone," then I consider that explicit and I expect it to continue to do so. (P4, Information Worker)

If an AI agent gets a 5 out of 5 on a customer query, it should do that again. (P3, AI Manager)

Memory Architecture & Lifecycle



Users want agent memory organized by work context, with labels and flexible ways to search and review

How should the agent's memories be organized?

How do users want agent memories to be structured and managed over time?

Re: organization, users most often described organizing memory around **projects, topics, workflows, case types, teams, clients/vendors, or business scenarios**, structures that reflect how work is actually carried out.

Participants also described other useful ways to classify memories, including **explicit vs. implicit**, or labels for **sensitivity or handling need**.

Many emphasized the need to **search, filter, and move across categories**, rather than rely on one fixed hierarchy, especially when memories cut across multiple topics.

A few also wanted a **date-based or chronological view**, as a useful way to review recency, trace decisions, or debug what led to an outcome.

Recommendation: Organize memories primarily around work contexts (projects, cases, workflows, clients, teams). Layer flexible labels and attributes on top (topic, explicit/implicit, sensitivity, handling). Support multiple ways to view the same memories, filtered views by labels, chronological view (by time) or searching. Filtered view (by labels)



I would want it organized by case type, like 'deposit change'... So, I would want to see by case type. (P3, AI Manager)

If an AI agent gets a 5 out of 5 on a customer query, it should do that again. (P3, AI Manager)



Users do not want a single default expiration period for memory

Should agent memories have a default time-to-live (an expiration period)?

Information Workers: prefer long-lived or user-controlled memory

Indefinite persistence reduces the productivity cost of re-teaching preferences, formats, and context. Many want memory deleted only when they explicitly initiate it, with optional deactivation of unused items to protect performance.

AI Managers: structured review cycle

Often favor 6-month reviews to account for policy updates, shifting business context, and changing data. Periodic reassessment helps avoid outdated, unsafe, or non-compliant agent behavior.

IT Admins: governance, cost, and regulatory retention drive lifespan

Expiration decisions may be constrained by storage budgets and mandates (e.g., multi-year retention in financial services). The priority is legal risk management and infrastructure impact, not individual preference.

Citizen Makers: need automated lifecycle management at scale

Manual memory management becomes burdensome. They advocated for intelligent handling such as archiving unused memories or condensing older items into higher-level summaries.

Shared conclusion: no universal default

Recommendation: Memory lifespan should be context-dependent and shaped by project duration, contractual terms, regulatory obligations, and how frequently or critically a memory influences agent behavior.

I don't think there is a perfect answer to this. Always it depends, on the severity and longevity of that specific decision. In certain events, no long-term persistent context. Sometimes there is a legal hold requirement, there are regulatory needs, so you have to store.
(P6, Citizen Maker)

It should stay there until I delete it. Otherwise, I'd have to provide more inputs.
(P11, Information Worker)

Memory in Collaborative and Shared Contexts

Shared memory should be layered and transparent, with clear boundaries to prevent leakage across people, projects, and contexts

Users want memory **separated into clear layers**: personal memory for individual preferences, shared project memory for common facts and decisions, and team-level memory for trusted practices others can rely on.

In collaborative settings, users **expect stronger transparency** into shared memory, what is being remembered, who added it, who validated it, when it was updated, and where it applies.

Shared memory is **most acceptable when it is tied to a common task**, feature or workflow.

The **biggest risk is leakage**: users worry that confidential, out-of-scope, or role-specific information could resurface for the wrong person, project, or task if boundaries are not explicit.

Recommendation: Design layered memory with explicit boundaries and visibility into what is shared, how it was learned, and where it applies.



So we introduced tiers. Not all memory is equal. 1 – individual. 2 – team memory but not shared, 3 – trusted memory, this is broadly applied and rarely questioned and used as a foundation for most of the automation. (P2, AI Manager)

It would need strict parameters... We work with a lot of confidential information; I would be concerned about how confidential info might be spat out later. (P4, Information Worker)



Shared agents should learn selectively, from validated, broadly useful patterns

Users are open to shared agents learning from other humans **when the system captures proven, broadly useful practices** rather than every individual behavior.

They especially **value agents learning from the most effective contributors**, validated team knowledge, or successful patterns that can help others work better.

At the same time, users **do not want shared agents to absorb inefficient work practices** or unproven ways of working that could reduce quality out habits, irrelevant noise, or for the group.

This makes **selective promotion important**: users expect systems to distinguish between observations, emerging patterns, and trusted knowledge before something shapes the broader team experience.

Recommendation: Enable agent memory to filter and promote only high-quality, validated patterns into shared learning



We have a human agent floor of 2500 human employees. We would not want an agent to learn accidentally from our lowest performers. We would want to put in some guardrails of only learning from cases with 4.5 stars or higher. If an agent is learning how to learn, we would want it learn from 5 - star cases.
(P3, AI Manager)

Sometimes the best way to learn is from other people... But maybe other people aren't efficient, so I don't want their process. You'd want the agent to learn from how everyone is using it and optimize for most efficient process. (P7, Information Worker)

Governance, Privacy & Risk

Users expect shared governance models with well defined role-based contributions for agent memory

Who should control agent memory? Should it vary by memory type or scenario?


Participants preferred **shared oversight across IT, legal/compliance and business leads**, with a strong desire for clear accountability to avoid inconsistent updates and maintain trust in the memory.

Some users described a **layered model** where end users contribute, designated owners review, and admins enforce rules or delete outdated memory.

One nuance to the expected shared governance model is that participants expect **individual ownership of contributed memories**, wanting control over what they create (e.g., ability to update/delete) without interference from others.

Multiple participants noted that **no established governance models** currently exist for managing agent memory within their organizations.

Recommendation: Design a memory management interface that enables a shared governance model with well defined structures (roles, processes, ownership boundaries).



No single person should own the agent memory. It should be like any other platform of shared knowledge. Ownership should follow the risk. (P2, AI Manager)

This is tough. Data and analytics team who handles nuts and bolts. There would need to be right authorizations. There is not a governance for this that would exist off hand. (P5, AI Manager)



Agent memory retaining or surfacing sensitive information in the wrong context is a core risk, requiring strict controls on storage, access, and retention

What are potential risks of agent memory and how can design or policy mitigate these?

What privacy safeguards and policies do users and admins expect?

Participants **emphasize “never store” categories** (PII, client, personal data, confidential data), driven by fear of leakage, cross-client exposure, or misuse. They expect to apply clear guardrails on what can never be stored for high-risk information.

Users want **tight access controls and segmentation** (role-based, geo-based, team-based), to prevent unauthorized cross-team or cross-client visibility.

There is consistent concern about **over-retention and lack of lifecycle control**, as memory policies must align to existing compliance, legal, and industry regulations. User expect retention policies, auto-deletion for compliance purposes, audit trails and ability to apply human review for compliance.

Recommendation: Design memory systems with clear rules for what cannot be stored, who can access data, and how long it is retained. Make these controls visible to users to build trust.



I'd want to impose rules, like don't remember anything of a personal nature. Don't remember or cross-reference agencies to potentially competing ones. Don't share cross-info back to clients. Don't mention one client to another or one department to another. (P13, IT Admin)

No client PII should ever stored. Do not remember this. You just verify this. The risk is that the agent could regurgitate that to another client. We've seen this with humans...we want to make that not storable. (P3, AI Manager)



Users expect agent-to-agent sharing to be selective, contextual, and tightly controlled

How should agent memory operate across multiple agents?

Should they share common memories/should memory transfer between tools?

Most participants **supported some memory transfer**, when continuity improves user experience (e.g., agent handoffs or orchestrated workflows).

Participants consistently **resist raw memory sharing across agents**, citing risks of error propagation, privacy leakage, and cross-context contamination.

Many prefer **controlled sharing of signals or outcomes** (not underlying memory), enabling agents to learn without exposing sensitive or irrelevant data.

Recommendation: Enable selective sharing between agents, where agents can exchange signals, summaries, or task-specific context (not raw memory), with clear user controls on what can be shared.



One thing that has evolved is how agents learning from each other. We have put in a structured learning process between the agents. Instead of sharing the memory, they learn from each other via signals that are purpose built. (P2, AI Manager)

In general, I don't think I have any issues with it [agents sharing memory]. I think it would be important to "teach them how to play" for what to share/not share, because we have some PII. Need some guard rails but in general it makes sense. (P5, AI Manager)



Participants want work and personal memory separate - except for communication style or details they explicitly choose to share


*If you use an AI agent at work, do you expect it to also know things about your personal life or preferences?
Why or why not?*

Most participants **defaulted to strict separation** of work and personal data, valuing AI's neutrality and worrying about indefinite retention and personal details being absorbed into company-owned systems.

Participants consistently **welcomed one exception: communication style**- tone, voice, and phrasing preferences, because it makes outputs sound authentic without exposing sensitive information.

Many expect **explicit user-controlled disclosure** (not autonomous gathering), with the ability to save personal details to memory only when they choose its beneficial (e.g., family details for employer sponsored benefits), and desire for incognito-style modes if personal memories could inadvertently be saved.

Recommendation: Default to a work-only memory scope, with opt-in personalization limited to communication style and clearly user-controlled disclosures.



This tool is as powerful as the data you give to it. It should only know if I feed that information into the tool. And then there's no going back. Keep it strictly for work. Because if I feed it to platform, then it's company's info. (P15, IT Admin)

I'm real big on communication style, when I'm reaching out to people and using a certain kind of script, so it's a blurred line because I want the agent to sound right. I tell ChatGPT it's too casual, or no rocket icons. It's teaching it "this is how I communicate." I want it to know my style. I don't think it really needs to know anything else. (P9, Citizen Maker)

Memory Models: Human vs. Agent



Users expect agent memory to outperform human memory on retention rigor, intentional forgetting and audibility.

Agent memory's value is precisely in capabilities human lack:


Superior factual retention and recall (e.g., dates, transactions, processes, client history) paired with fast recall that humans can't match.

Cross-context connections that humans may miss when narrowly focused (e.g., supply chain ↔ demand planning, prior client interactions surfaced before a 1:1). The ability to recall information in relevant contexts.

Intentional forgetting: Critically, users also expect agent forgetting to be *better* than human forgetting. Humans forget randomly and lose important things, while the agent should decay low-value patterns intentionally, on visible rules, with user controls to pin or remove items.

Auditability, which users repeatedly invoke as the defining contrast to human memory: where human memory is automatic and fuzzy, agent memory should be inspectable, explainable, and bounded. Users want to see what's been retained, why, and remove it.

Recommendation: Build memory around rigorous, auditable factual recall. Treat forgetting as an explicit feature with visible decay logic and user controls.



The value comes from it not being a human and doing things humans aren't capable of....it has the benefit of more rigorous retention, which we are not capable of. It can be more structured and maintain boundaries vs. humans being more fluid. (P4, Information Worker)

Where it should be different, it should be auditable. It gives me greater confidence. It should forget intentionally, it should only decay low value patterns. Humans forget randomly. (P2)



Users expect human memory to remain superior on emotional inference and relational discretion


Participants are clear that some forms of memory belong with humans, not the agent.

Emotional inference reading tone, morale, frustration, humor is something users want the agent to perceive in the moment but never retain or use as a decision input; multiple participants drew an explicit line between "informing my judgment" (acceptable) and "making decisions on tone" (not acceptable).

Relational discretion in what to store and what to recall. This is the strongest do-not-store boundary: client personal details (health, family), colleague disclosures, team workload and morale, off-hand personal context — users worry about both privacy and over-personalization.

Holistic judgment is the third gap: humans synthesize across domains and lived experience in ways the agent can't, since it only knows what it has been fed; users want to keep that integrative role for themselves rather than have the agent simulate it.

Recommendation: Default-exclude emotional, relational, and personal content from durable memory.



Humans have more short-term memory, long-term memory, mid-term is grey. I think for humans this is okay. I wouldn't want an agent to have super focused short-term memory, forget mid-term and then a legacy memory of one time I got a 5 star a review. (P3)

Where it should be different, it should be auditable. It gives me greater confidence. It should forget intentionally, it should only decay low value patterns. Humans forget randomly. (P2)

By Persona



Memory Expectations Differ by Persona



IT Admins

Governed resource.
Balance usability with storage cost, security, legal exposure, and organizational hierarchy.



AI Managers

Decision-input system.
Need auditability, compliance alignment, and structured reviews; emphasize confidence and impact.



Citizen Makers

Scalable infrastructure.
Push automation, condensation, and proactive context surfacing over manual curation.



Information Workers

Personal productivity asset.
Prioritize stability, low friction, and clear controls to add, update, or delete—no internals required.

4

Design Opportunities

Design Opportunities

Six priorities for enterprise agent memory, grounded in user expectations

01

Hybrid Learning, with user in control

Implicit learning for routine work, explicit for high-stakes rules, suggest-to-save in between with user able to validate what is promoted.

02

Organize Memory Around Work Context

Anchor memory to projects and workflows; layer flexible views by topic, time, sensitivity.

03

Lifecycle by Context, Not by Default

Lifecycle rules are tied to projects, contracts, and regulation, with no universal expiration period for memories.

04

Layered, Transparent Sharing in Collaborative Contexts

Personal → project → trusted tiers with provenance and leakage guardrails to prevent sharing of personal memories in group contexts.

05

Shared Governance, with clearly defined ownership

Role-based oversight; sensitive categories blocked from storage by default

06

Auditable when needed

Inspectable memory, visible decay, persona-tuned surfaces; transparent when needed

Appendix

Next Steps from AI Memory Top 8 Emerging Insights

Section	Purpose	Key Next Steps
 Amplify & Align In progress	Align internal teams by broadly sharing these emerging research insights for awareness. Solicit open research questions for future exploration.	• Share emerging findings widely (partner teams, forums) to ensure everyone is on the same page. • Collect feedback on open research questions.
 Generative Research In progress	Deepen our understanding of user expectations, mental models, and concerns about agent memory, to guide future designs and feature capabilities.	• Investigate user mental models of “memory” (what users expect AI to remember and how they want to manage it). • Examine privacy expectations and trust boundaries; determine users’ preferred control mechanisms for memory in collaborative contexts.
 Competitive Landscape Synthesis Future	Stay informed on competitors’ advances in AI memory to identify strategic opportunities for Microsoft’s differentiation.	• Monitor evolving memory features and strategies from key competitors (OpenAI, Anthropic, Google). • Identify emerging norms in the market and areas where Microsoft can lead or offer unique value.
 Design Evaluations Future	Ensure upcoming memory design solutions are user-tested and meet essential criteria (usability, transparency, control) before launch.	• Plan usability testing and scenario-based evaluations for memory features once designs are available. • Validate that new memory features for enterprise specific personas provide adequate transparency, user control, and contextual relevance in real workflows.
 Metrics & Evaluation Framework Future	Establish how success will be measured and how strategic impact should be demonstrated over time.	• Define clear success metrics for agent memory (e.g., user trust, efficiency, accuracy). • Implement instrumentation and evaluation methods to track these metrics and memory’s impact over time.

Preferred Learning Mode by Task Ambiguity & Risk

Preference shifts by task type — not by a single fixed user stance

